

How Amazon S3 Works ✨

#59: Break Into Amazon Engineering (4 Minutes)



NEO KIM

OCT 25, 2024



193



6



17

Get my system design playbook for FREE on newsletter signup:

Subscribe

This post outlines the internal architecture of AWS S3. You will find references at the bottom of this page if you want to go deeper.

- [Share this post](#) & I'll send you some rewards for the referrals.

Note: This post is based on my research and may differ from real-world implementation.

Once upon a time, there was an analytics startup.

They collect data from customer websites and store it in log files.

Yet they had only a few customers.

So a tiny storage server was enough.

But one morning they got a new customer with an extremely popular website.

And number of log files started to skyrocket.

Yet their storage server had only limited capacity.



So they bought new hardware.

Although it temporarily solved their storage issues, there were newer problems.

Here are some of them:

1. Scalability

The storage server might become a capacity bottleneck over time.

While installation and maintenance of a larger storage server is expensive.

2. Performance

The storage server must be optimized for performance.

But they didn't have the time and knowledge for it.

Onward.

[Product for Engineers - Sponsor](https://newsletter.systemdesign.one/p/s3-architecture)



Product for Engineers is PostHog's newsletter dedicated to helping engineers improve their product skills. It features curated advice on building great products, lessons learned (and mistakes) from building PostHog, and research into the practices of top startups.

They wanted to ditch the storage management problem.

And focus only on product development.

So they moved to Amazon Simple Storage Service (**S3**) - an object storage.

It stores [unstructured data](#) without hierarchy.

And handle 100 million requests per second.


Yet having performance at scale is a hard problem.

So smart engineers at Amazon used simple ideas to solve it.

S3 Architecture

Here's how S3 works:

1. Scalability



They provide REST API via the web server.

While metadata & file content are stored separately - it lets them scale easily.

They store the metadata of uploaded data objects in a key-value database. And cache it for high availability.


Each component in the above diagram consists of many microservices. While services interact with each other via [API contracts](#).

Ready for the best part?

2. Performance

They store uploaded data in mechanical hard disks to reduce costs.

And organize data on the hard disk using [ShardStore](#) - it gives better performance. Think of **ShardStore** as a variant of log-structured merge ([LSM](#)) tree data structure.



newsletter.systemdesign.one

Moving Parts of the Hard Disk


A larger hard disk can store more data.

But seek & rotation times remain constant due to moving parts. So its throughput is the same as a small disk.

Put simply, a larger disk performs poorly in retrieving data.

- **Throughput** means the amount of data transferred over time - measured in MB/s
- Imagine the **seek time** as time needed to move the head to a specific track on the disk.
- Think of **rotation time** as the time needed for the head to reach a specific point of data.

Also a single disk might become a hot spot if the data isn't distributed uniformly across disks.




Reading Data in Parallel for High Throughput

So they *replicate data* across many disks and do *parallel reads* - it gives higher throughput.

Besides the load on a specific disk is lower because data can be read from any di
Thus preventing hot spots.

Yet full data replication is expensive from a storage perspective.



Replicating Data With Erasure Coding


So they use [erasure coding](https://newsletter.systemdesign.one/p/s3-architecture) to replicate data.

Think of **erasure coding** as a technique to replicate data with a smaller storage overhead.

Here's how it works:

- A data object is split into pieces called **identity shards**.
- Mathematical algorithms are used to create extra chunks called **parity shards**.
- The number of parity shards is lower than identity shards.

The parity shards contain enough information to recreate any identity shards. The erasure coding offers the same level of fault tolerance as full replication.




Recreating the Data Object From Identity and Parity Shards

Also any combination of identity and parity shards can recreate the data object. So there's no need to replicate the entire dataset, thus reducing storage needs.

Besides they store shards across different hard disks to avoid hot spots. And data objects from a single customer are distributed across disks for performance.

systemdesignone



S3 stores 280 trillion data objects across a million hard disks - zettabytes of data.

And offer 99.99% availability.

This case study shows you don't need expensive hardware to scale with performance.

Subscribe to get simplified case studies delivered straight to your inbox:

Type your email...	Subscribe
--------------------	-----------



Follow me on [LinkedIn](#) | [YouTube](#) | [Threads](#) | [Twitter](#) | [Instagram](#)

Thank you for supporting this newsletter. Consider sharing this post with your friends and get rewards. Y'all are the best.


1 Referral  Leetcode Master Template	2 Referrals  Interview Mistakes to Avoid PDF	3 Referrals  Popular Interview Questions PDF
---	--	---



How Google Search Works 🔥

NEO KIM • 30 SEPTEMBER 2024

[Read full story →](#)



How Facebook Was Able to Support a Billion Users via Software Balancer ⚡

NEO KIM • 11 OCTOBER 2024

[Read full story →](#)

References

- [FAST '23 - Building and Operating a Pretty Big Storage System \(My Adventure with Amazon S3\)](#).
- [Building and operating a pretty big storage system called S3](#)
- [AWS re: Invent 2023 - Dive deep on Amazon S3 \(STG314\)](#).

- [Using Lightweight Formal Methods to Validate a Key-Value Storage Node in Amazon S3](#)
- [AWS is building a new storage backend for S3 on 40k lines of Rust](#)
- [How a Hard Drive Works](#)
- [Understanding Erasure Coding](#)
- [What is Flash Storage?](#)
- [How is disk speed measured and what is fast?](#)
- Block diagrams created with [Eraser](#)

Subscribe to The System Design Newsletter

By Neo Kim · Hundreds of paid subscribers

Download my system design playbook on newsletter signup for FREE

By subscribing, I agree to Substack's [Terms of Use](#), and acknowledge its [Information Collection Notice](#) and [Privacy Policy](#).



193 Likes · 17 Restacks

← Previous

Next →

Discussion about this post

Comments

Restacks



Write a comment...



Manpreet Singh 26 Oct 2024

♥ Liked by Neo Kim

Great research, I have shared it via email to my friends.

♥ LIKE (2) 💬 REPLY

2 replies by Neo Kim and others



Susmit Sarkar 26 Oct 2024

♥ Liked by Neo Kim

Good writeup, but please do read Alex Xus write up as well in ByteByte Go !

Not comparing both as it would be unfair

♥ LIKE (1) 💬 REPLY

2 replies by Neo Kim and others

4 more comments...

© 2025 Neo Kim · [Publisher Privacy](#)

[Substack](#) · [Privacy](#) · [Terms](#) · [Collection notice](#)

[Substack](#) is the home for great culture